

Ruprecht Karls University Heidelberg
Institute of Computer Science
Database Systems Research Group

Bachelor Thesis
Supporting Offline Editing in Mobile and
Collaborative Apps

Name: Mirko Kiefer

I declare that this thesis was composed by myself and that the work contained therein is my own, except where explicitly stated otherwise in the text.

Date of Submission: 17. April 2013

Abstract

Zusammenfassung

Inhaltsverzeichnis

1	Introduction	1
1.1	Motivation	1
1.2	Goals of the thesis	2
1.3	Structure of the thesis	2
2	Background	3
2.1	Causality Preservation	3
2.2	Optimistic Synchronization	3
2.3	Vector Clocks	3
2.4	Operational Transformation	3
2.5	Commutative Replicated Data Types	3
2.6	State vs. Edit-Based Syncing	3
2.7	Three-Way Merging	3
2.8	Most Recent Common Ancestor	3
3	Designing a Synchronization Framework	4
3.1	Application Scenarios	4
3.1.1	Relational Data Synchronization	4
3.1.2	Hierarchical Data Synchronization	5
3.1.3	Text Synchronization	5
3.2	Requirements	5
3.3	Existing Systems	6
3.3.1	git	6
3.3.2	CouchDB	6
3.4	Evaluation of Existing Systems	6
3.5	Architecture of Synclib	7
3.6	Implementation	7
3.6.1	Technology	7
3.6.2	Data Model	8
3.6.3	Change History Tracking	8

3.6.4	Most Recent Common Ancestor	8
3.6.5	Synchronization Protocol	8
3.6.6	Conflict Handling	9
3.6.7	Integration with existing Application Logic	9
4	Evaluation	10
	Literaturverzeichnis	11

1 Introduction

1.1 Motivation

Applications that allow users to collaborate on data on a central server are in widespread use. Popular examples are document authoring tools like Google Docs, project collaboration apps like Basecamp or Trello or even large scale collaboration projects like Wikipedia.

The traditional architecture of these applications follows a client-server model where the server hosts the entire application logic and persistence. Users access the application through a thin client, most commonly a web browser, who only runs the logic required for rendering the user interface.

This model works well when using desktop computers with a reliable, high-speed connection to the server.

The widespread adoption of mobile devices forces developers to re-think this architecture. Users can now carry their devices with them and expect their applications to work outside their home or office network. Applications therefore have to work with limited mobile internet access or often no access at all.

The only way to support this is by moving relevant parts of the application logic to the client and by replicating data for offline use.

The new architecture comes at a high price - the additional client logic and persistence adds a lot of complexity. While in the thin client-fat server model developers only had to maintain a single technology set on the server, they now face different technologies on each platform they aim to support with a fat client.

In addition to that the ability to edit data offline requires an entire new layer of application logic to manage the propagation and merging of changes and to resolve conflicts.

- users have multiple devices - only partial data required offline

1.2 Goals of the thesis

Thesis problem statement This thesis aims to develop patterns and tools to make the development of offline capable, collaborative apps more productive and making their maintenance manageable.

The focus lies on two core topics:

- Managing distributed application logic
- Synchronizing offline changes

These topics are closely related - if an application has to be available offline, the developer has to build client logic necessary to replace the server's tasks. Part of that logic is the ability to synchronize data with the server or even directly with other clients. The additional client logic has to be maintained and upgraded for new releases of the application. As the client logic is distributed among all users of the application, a code upgrade becomes more complex to manage than a simple server update. We will see how the same logic used to synchronize application data can be used for updating distributed application code.

- 3-way merging for ordered XML is solved - focus on interfaces / integration into existing code - composite data structures / synchronization granularity - peer2peer (synchronization topology) - how manage constraints?

1.3 Structure of the thesis

Here you describe the structure of the thesis. For example:

In Kapitel 2 werden grundlegende Methoden für diese Arbeit vorgestellt.

2 Background

Here you discuss some basics for your work and outline existing research in the area of your thesis by citing research papers like [1] by Lindholm and [2, 3] Candia.

2.1 Causality Preservation

2.2 Optimistic Synchronization

2.3 Vector Clocks

2.4 Operational Transformation

2.5 Commutative Replicated Data Types

2.6 State vs. Edit-Based Syncing

2.7 Three-Way Merging

- 3DM tool (Lindholm)

2.8 Most Recent Common Ancestor

3 Designing a Synchronization Framework

3.1 Application Scenarios

We describe common synchronization scenarios based on popular mobile applications.

3.1.1 Relational Data Synchronization

The Wunderlist app serves as an example for a common kind of data model that requires syncing of data in a relational schema.

Wunderlist's schema could be defined as the following:

- User (name, email, has Todo Lists)
- Invited User (name, email)
- Invited User List (has Invited Users)
- Todo Item (title, description, due date, belongs to Todo List)
- Todo List (name, belongs to Users, has Todo Items)

The User type has a singleton instance who represents the user of the app.

Users can be invited to Todo Lists. As their list of Todo Lists is hidden from the current user Invited User is a separate type.

Invited User List is simply a cached list of all users that have been invited in the past. While Invited User List is an unordered list, Todo Lists and Todo Items are ordered.

Syncing lists of unordered object IDs never causes conflicts while syncing ordered object IDs can cause order conflicts.

3.1.2 Hierarchical Data Synchronization

Dropbox synchronizes a file system - it is therefore a good example for syncing of hierarchical data.

The data model is simple:

- Tree Item (name)
- Tree extends Tree Item (has children of type Tree Items)
- Data extends Tree Item (data)

The list of child Tree Items can either be ordered or unordered. While Dropbox does not sync the order of files there are scenarios where this is required.

Syncing trees can trigger conflicts if sub trees have been modified concurrently.

3.1.3 Text Synchronization

Collaborative document editors like Google Docs need to synchronize text that is concurrently edited.

Google Docs currently does not support offline editing.

Syncing text is equal to the problem of syncing an ordered list and can trigger conflicts.

3.2 Requirements

From the common scenarios we derive a set of requirements for a synchronization solution.

Requirements for strategies:

- Causality preservation
- Eventual consistency
- Optimistic synchronization
- Expose conflicts
- Support peer-to-peer or hybrid synchronization
- Integration with existing app logic

(TODO: need to explain why this set of requirements, constraints on mobile devices...)

Aspects to consider when evaluating strategies:

- How are updates detected?
- How are updates propagated? (Stream or Snapshot)
- How are updates merged/reconciled? (State or Edit-based)
- Level of structural awareness (Textual, Syntactic, Semantic/Structural)

3.3 Existing Systems

3.3.1 git

- Data structure: filesystem/tree
- Merging: tree-based, three-way merge
- Propagation: snapshot-based
- Supports peer-to-peer

3.3.2 CouchDB

- Data structure: key-value
- Merging: tree-based
- Propagation: stream-based
- Supports peer-to-peer

3.4 Evaluation of Existing Systems

Here we evaluate solutions like CouchDB, Dropbox, iCloud, git, (Google Docs), (Parse.com) based on the requirements.

3.5 Architecture of Synclib

Based on the requirements and the evaluation of existing systems we derive a unique architecture for a practical synchronization solution.

- **no timestamps:** state-based 3-way merging
- **no change tracing:** change tracing is not necessary - support diff computation on the fly
- **data agnostic:** leave diff and merge of the actual data to plugins
- **distributed:** syncing does not require a central server
- **be small:** only implement the functional parts of syncing - leave everything else to the application (transport, persistence)
- **sensitive defaults:** have defaults that *just work* but still support custom logic (e.g. for conflict resolution)

3.6 Implementation

- cross-platform through web standards - solve server behaviour through native proxy
- diff-merge-patch - most-recent-common-ancestor - focus on ordered-sets with moves
(Todo list)

We describe implementation details like the technologies used, code structure and the testing framework to evaluate the system.

As syncing is state based we need to track the history of edits on each client. Each client has his own replica of the database and commits data locally. On every commit we create a commit object that links both to the new version of the data and the previous commit. If a client is connected to a server he will start the sync process on every commit. As synclib2's architecture is distributed a server could itself be a client who is connected to other servers.

To the latest commit on a database we refer to as the 'head'.

3.6.1 Technology

- everything web-based -> only way to be cross-platform - note on alternatives (Lua, native)

3.6.2 Data Model

- support tracking different data structures (sets, lists, ordered sets and lists, dictionaries, trees, composite data structures) - need to implement diff, merge, patch for each - ensure conflicts are made explicit

3.6.3 Change History Tracking

- use content-addressable store - only store changes and reference unchanged data through hashes -> like git - commit links to data and parent commit

3.6.4 Most Recent Common Ancestor

- implementation as separate module - describe algorithm in background

3.6.5 Synchronization Protocol

Syncing follows the following protocol:

Client has committed to its local database.

Client pushes all commits since the last synced commit to Server.

Client asks Server for the common ancestor of client's head and the server's head

Client pushes all changed data since the common ancestor to Server.

```
if common ancestor == server head
    // there is no data to merge
    try fast-forward of server's head to client's head
    if failed (someone else updated server's head in the meantime) then start over
else
    Client asks Server for all commits + data since the common ancestor
    Client does a local merge and commits it to the local database
    start over
```

This protocol is able to minimize the amount of data sent between synced stores even in a distributed, peer-to-peer setting.

Updating the server's head uses optimistic locking. To update the head you need to include the last read head in your request.

3.6.6 Conflict Handling

3.6.7 Integration with existing Application Logic

- demonstrate how to interface with standard MVC frameworks like Backbone, Ember.js

4 Evaluation

We evaluate the implementation based on the set of requirements specified in section 3.2.

Evaluate the proof-of-concept by simulating syncing of data structures used in the problem scenarios with realistic network latency and disconnection.

show efficiency both on client-server and peer2peer.

Literaturverzeichnis

- [1] Tancred Lindholm. XML-aware data synchronization for mobile devices. 2009.
- [2] Giuseppe DeCandia, Deniz Hastorun, Madan Jampani, Gunavardhan Kakulapati, Avinash Lakshman, Alex Pilchin, Swaminathan Sivasubramanian, Peter Voshall, and Werner Vogels. Dynamo: amazon’s highly available key-value store. 41(6):205–220, 2007.
- [3] David Ratner, Peter Reiher, Gerald J Popek, and Geoffrey H Kuenning. Replication requirements in mobile environments. *Mobile Networks and Applications*, 6(6):525–533, 2001.